

Lab 8: Tree-based Models

We will use the `Carseats` data set in the `ISLR` package to predict `high_sales` for carseats at 400 different stores.

```
library(ISLR) ## data package
library(tidyverse) ## data manipulation
library(tidymodels) ## tidy modeling
library(knitr) ## tables

## reproducible
set.seed(445)

## data
str(Carseats)

## 'data.frame':   400 obs. of  11 variables:
## $ Sales       : num  9.5 11.22 10.06 7.4 4.15 ...
## $ CompPrice   : num  138 111 113 117 141 124 115 136 132 132 ...
## $ Income      : num   73 48 35 100 64 113 105 81 110 113 ...
## $ Advertising: num   11 16 10 4 3 13 0 15 0 0 ...
## $ Population  : num  276 260 269 466 340 501 45 425 108 131 ...
## $ Price       : num  120 83 80 97 128 72 108 120 124 124 ...
## $ ShelveLoc   : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1
1 3 2 3 3 ...
## $ Age        : num   42 65 59 55 38 78 71 67 76 76 ...
## $ Education   : num   17 10 12 14 13 16 15 10 10 17 ...
## $ Urban       : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1
...
## $ US         : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2
...
```

0.1 Data Preparation

1. Make a copy of the `Carseats` data frame called `df`.
2. Create a variable called `high_sales` in `df` that takes the value “high” if `Sales > 8` and “low” otherwise.
3. Convert your `high_sales` column to be a factor.
4. Remove the `Sales` column from `df`.

0.2 Decision Trees

The `rpart.plot` package can be useful for creating tree diagrams.

```
library(rpart.plot) ## plotting trees
```

1. Using the `decision_tree()` function with the “`rpart`” engine in `tidymodels`, fit a large classification tree to predict `high_sales` using every variable in `df`. [**Hint:** The syntax is very similar fitting a `linear_reg`]
2. Inspect your tree. How many terminal nodes do you have? What is the training error rate?
3. Use the `rpart.plot` function to visualize your tree. What is the most important indicator of high sales?

[**Hint:** You need to extract the fit from your fitted tree using `extract_fit_engine()` before plotting.]

4. Split your observations into a training and a test set with 200 records each. Estimate the test error rate of your tree.
5. Produce a confusion matrix for your test set.
6. Perform cross-validation to determine the optimal level of tree complexity on your training data set. Which α (corresponds to `k` in the output) should we choose?
7. Use the functions `finalize_workflow` and `select_best()` to prune your tree to the chosen complexity. Plot your final tree.
8. Repeat 4-5 using your pruned tree. Which performs better?

0.3 Bagging & Random Forests

We will use the `rand_forest` function to perform bagging and random forests. Recall that bagging is simply a special case of random forests with $m = p$. Here is an example of a bagging specification for classification:

```
bagging_spec <- rand_forest(mtry = .cols()) |> ## automatically grabs
  the number of columns
  set_engine("randomForest", importance = TRUE) |> ## save information
  to plot importance
  set_mode("regression")
```

The `vip` package can be used to easily plot variable importance.

```
library(vip)

##
## Attaching package: 'vip'

## The following object is masked from 'package:utils':
##
##      vi
```

1. Perform bagging on your training `df` to predict `high_sales`. Specify `importance = TRUE` to also obtain information on the importance of each predictor.
2. Make a plot of the importance values for each predictor using the `vip` function. What is the predictor with the highest importance?
3. Estimate the test error rate using your bagged tree model.
4. Repeat 1-3 using a random forest with $m = \sqrt{p}$ via `mtry = sqrt(.cols())`.
5. Compare the OOB confusion matrix to your test confusion matrix. [**Hint:** The `confusion` element of the model output `fit` is OOB.]

0.4 Boosting

To perform boosting we will use the `boost_tree()` function with the "xgboost" engine.

Here is an example specification:

```
boost_spec <- boost_tree(trees = 5000, tree_depth = 4) %>%  
  set_engine("xgboost") %>%  
  set_mode("classification")
```

1. Fit a boosted tree ensemble to your training `df` predicting `high_sales` with $B = 5,000$ trees, learning rate of $\lambda = 0.01$, and an interaction depth of $d = 2$.
2. Estimate the test error rate using your boosted tree model and compare to all previously fit models.